

✓ Creating DataFrames

We will create a `DataFrame` object from a dataset in a CSV file using the `read_csv` method on the pandas - `pd` - object. While the `read_csv` method can read dataset directly from a web URL, it's best to download to your computer first to ensure the data is persisted.

About the Data

We will use the dataset on nobel laureates available at: <http://api.nobelprize.org/v1/laureates.csv>. Download the file to your computer using: `curl -OL http://api.nobelprize.org/v1/laureates.csv`

Imports

```
import datetime
import numpy as np
import pandas as pd
```

✓ Creating a Series

```
np.random.seed(0) # set a seed for reproducibility
pd.Series(np.random.rand(5), name='random')
```

```
0    0.548814
1    0.715189
2    0.602763
3    0.544883
4    0.423655
Name: random, dtype: float64
```

✓ Creating a DataFrame from a Series

Use the `to_frame()` method:

```
pd.Series(np.linspace(0, 10, num=5)).to_frame()
```

| | 0 | |
|---|------|--|
| 0 | 0.0 | |
| 1 | 2.5 | |
| 2 | 5.0 | |
| 3 | 7.5 | |
| 4 | 10.0 | |

✓ From a list of dictionaries

```
pd.DataFrame([
    {'mag' : 5.2, 'place' : 'California'},
    {'mag' : 1.2, 'place' : 'Alaska'},
    {'mag' : 0.2, 'place' : 'California'},
])
```

| | mag | place | |
|---|-----|------------|--|
| 0 | 5.2 | California | |
| 1 | 1.2 | Alaska | |
| 2 | 0.2 | California | |

✓ From a NumPy array

```
pd.DataFrame(
    np.array([
        [0, 0, 0],
        [1, 1, 1],
        [2, 4, 8],
        [3, 9, 27],
        [4, 16, 64]
    ]), columns=['n', 'n_squared', 'n_cubed']
)
```

| | n | n_squared | n_cubed |
|---|---|-----------|---------|
| 0 | 0 | 0 | 0 |
| 1 | 1 | 1 | 1 |
| 2 | 2 | 4 | 8 |
| 3 | 3 | 9 | 27 |
| 4 | 4 | 16 | 64 |

✓ Creating a DataFrame by Reading in a CSV File

Finding information on the file before reading it in

Before attempting to read in a file, we can use the command line to see important information about the file that may determine how we read it in. We can run command line code from Jupyter Notebooks (thanks to IPython) by using `!` before the code.

Number of lines (row count)

For example, we can find out how many lines are in the file by using the `wc` utility (word count) and counting lines in the file (`-l`). Run the cell below to confirm the file has 1002 lines:

```
!wc -l laureates.csv # this will not work on Windows commandline
# On windows the !dir command will show directory contents

1002 laureates.csv
```

We can even capture the result of a command and use it in our Python code:

✓ Reading in the file

Our file is small in size, has headers in the first row, and is comma-separated, so we don't need to provide any additional arguments to read in the file with `pd.read_csv()`, but be sure to check the [documentation](#) for possible arguments. To read data from file we can use `pd.read_csv()` and for other delimited files, such as tab (`\t`), we can use the `read_csv()` function with the `sep` argument equal to the delimiter. We can use the `read_excel()` function for Excel files, the `read_json()` function for JSON (JavaScript Object Notation) files

```
import pandas as pd
df = pd.read_csv('laureates.csv')
```

✓ Getting Documentation on Python elements

You can utilize the built-in `help()` function for documentation on Python elements. Simply run `help()`, passing in the package, module, class, object, method, or function. Assuming we aliased pandas as `pd` when we imported it, we can run `help(pd)` to see information on the pandas package; `help(pd.DataFrame)` for all the methods and attributes of a dataframe (note we can also pass in an already created DataFrame object instead); and `help(pd.read_csv)` to learn more about the pandas function for reading CSV files into Python and how to use it

Run the code cell below to see documentation for the DataFrame class.

```
import pandas as pd
help(pd.DataFrame)
```

```

    __lt__(self, other)
        Return self<value.
    __mod__(self, other)
    __mul__(self, other)
    __ne__(self, other)
        Return self!=value.
    __or__(self, other)
        Return self|value.
    __pow__(self, other)
    __radd__(self, other)
    __rand__(self, other)
    __rfloordiv__(self, other)
    __rmod__(self, other)
    __rmul__(self, other)
    __ror__(self, other)
        Return value|self.
    __rpow__(self, other)
    __rsub__(self, other)
    __rtruediv__(self, other)
    __rxor__(self, other)
    __sub__(self, other)
    __truediv__(self, other)
    __xor__(self, other)
    -----
    Data and other attributes inherited from pandas.core.arraylike.OpsMixin:
    __hash__ = None

```

Let's review summary statistics for the nobel laureates dataframe - df: use `describe()` .

```
df.describe()
```

`df.describe()` does not really tell us much. The `info()` method provides more information as you can check by running the code cell below.

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    1000 non-null  int64
 1   firstname             1000 non-null  object
 2   surname               968 non-null   object
 3   born                  999 non-null   object
 4   died                  1000 non-null  object
 5   bornCountry           969 non-null   object
 6   bornCountryCode       969 non-null   object
 7   bornCity              966 non-null   object
 8   diedCountry           653 non-null   object
 9   diedCountryCode       653 non-null   object
10   diedCity              647 non-null   object
11   gender                1000 non-null  object
12   year                  1000 non-null  int64
13   category              1000 non-null  object
14   overallMotivation     23 non-null    object
15   share                 1000 non-null  int64
16   motivation            1000 non-null  object
17   name                  736 non-null   object
18   city                  731 non-null   object
19   country               733 non-null   object
dtypes: int64(3), object(17)
memory usage: 156.4+ KB

```

We can use the dataframe `head()` or `tail()` method to view some actual entries. Without a numeric parameter, both methods return 5 entries!

Run the next 2 code cells below to see the first 10 and the last 10 entries

```
df.head(10)
```

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity |
|---|----|----------------|-----------|------------|------------|-----------------------------|-----------------|-----------------------------|
| 0 | 1 | Wilhelm Conrad | Röntgen | 1845-03-27 | 1923-02-10 | Prussia (now Germany) | DE | Lennep (now Remscheid) |
| 1 | 2 | Hendrik A. | Lorentz | 1853-07-18 | 1928-02-04 | the Netherlands | NL | Arnhem |
| 2 | 3 | Pieter | Zeeman | 1865-05-25 | 1943-10-09 | the Netherlands | NL | Zonnemaire |
| 3 | 4 | Henri | Becquerel | 1852-12-15 | 1908-08-25 | France | FR | Paris |
| 4 | 5 | Pierre | Curie | 1859-05-15 | 1906-04-19 | France | FR | Paris |
| 5 | 6 | Marie | Curie | 1867-11-07 | 1934-07-04 | Russian Empire (now Poland) | PL | Warsaw |
| 6 | 6 | Marie | Curie | 1867-11-07 | 1934-07-04 | Russian Empire (now Poland) | PL | Warsaw |
| 7 | 8 | Lord | Rayleigh | 1842-11-12 | 1919-06-30 | United Kingdom | GB | Langford Grove Maldon Essex |
| 8 | 9 | Philipp | Lenard | 1862-06-07 | 1947-05-20 | Hungary (now Slovakia) | SK | Pressburg (now Bratislava) |
| 9 | 10 | J.J. | Thomson | 1856-12-18 | 1940-08-30 | United Kingdom | GB | Cheetham Hill |

Next steps:

Generate code with df

 View recommended plots

```
df.tail(10)
```

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | borr |
|-----|------|-----------|------------|------------|------------|--|-----------------|------|
| 990 | 1025 | Drew | Weissman | 1959-09-07 | 0000-00-00 | USA | US | Lex |
| 991 | 1026 | Pierre | Agostini | 1941-07-23 | 0000-00-00 | French protectorate of Tunisia (now Tunisia) | TN | |
| 992 | 1027 | Ferenc | Krausz | 1962-05-17 | 0000-00-00 | Hungary | HU | |
| 993 | 1028 | Anne | L'Huillier | 1958-08-16 | 0000-00-00 | France | FR | |
| 994 | 1029 | Moungi | Bawendi | 1961-00-00 | 0000-00-00 | France | FR | |
| 995 | 1030 | Louis | Brus | 1943-00-00 | 0000-00-00 | USA | US | Clev |
| 996 | 1031 | Aleksey | Yekimov | 1945-00-00 | 0000-00-00 | USSR (now Russia) | RU | |
| 997 | 1032 | Jon | Fosse | 1959-09-29 | 0000-00-00 | Norway | NO | Haug |
| 998 | 1033 | Narges | Mohammadi | 1972-04-21 | 0000-00-00 | Iran | IR | Z |
| 999 | 1034 | Claudia | Goldin | 1946-00-00 | 0000-00-00 | USA | US | Nev |

Both `head()` and `tail()` and indeed `select` statements will return the full set of attributes for the entries. How can we project to just a select set of attributes? This will be the equivalent of an SQL project operator listing the output columns.

Run the code cell below to show only 'firstname','surname','city' and 'country' columns for the first 10 entries.

```
df.head(10)[['firstname','surname','city','country']]
```

| | firstname | surname | city | country | |
|---|----------------|-----------|-----------|-----------------|--|
| 0 | Wilhelm Conrad | Röntgen | Munich | Germany | |
| 1 | Hendrik A. | Lorentz | Leiden | the Netherlands | |
| 2 | Pieter | Zeeman | Amsterdam | the Netherlands | |
| 3 | Henri | Becquerel | Paris | France | |
| 4 | Pierre | Curie | Paris | France | |
| 5 | Marie | Curie | NaN | NaN | |
| 6 | Marie | Curie | Paris | France | |
| 7 | Lord | Rayleigh | London | United Kingdom | |
| 8 | Philipp | Lenard | Kiel | Germany | |
| 9 | J.J. | Thomson | Cambridge | United Kingdom | |

Next steps: [Generate code with df](#) [View recommended plots](#)

Following the example of the last code cell, write a code cell below to return the last 10 records of nobel laureates but only showing their firstname, surname, year of birth, year of death and country

```
# Add your code below this line
df.head(10)[['firstname', 'surname', 'born', 'died', 'country']]
```

| | firstname | surname | born | died | country | |
|---|----------------|-----------|------------|------------|-----------------|--|
| 0 | Wilhelm Conrad | Röntgen | 1845-03-27 | 1923-02-10 | Germany | |
| 1 | Hendrik A. | Lorentz | 1853-07-18 | 1928-02-04 | the Netherlands | |
| 2 | Pieter | Zeeman | 1865-05-25 | 1943-10-09 | the Netherlands | |
| 3 | Henri | Becquerel | 1852-12-15 | 1908-08-25 | France | |
| 4 | Pierre | Curie | 1859-05-15 | 1906-04-19 | France | |
| 5 | Marie | Curie | 1867-11-07 | 1934-07-04 | NaN | |
| 6 | Marie | Curie | 1867-11-07 | 1934-07-04 | France | |
| 7 | Lord | Rayleigh | 1842-11-12 | 1919-06-30 | United Kingdom | |
| 8 | Philipp | Lenard | 1862-06-07 | 1947-05-20 | Germany | |
| 9 | J.J. | Thomson | 1856-12-18 | 1940-08-30 | United Kingdom | |

Next steps: [Generate code with df](#) [View recommended plots](#)

Querying & Locating Data in the DataFrame

One of the most useful tasks in pandas is locating data that satisfies desired criteria. For example, we can locate a Nobel laureate with a particular surname. Let’s look at the record of Caltech’s most beloved figure, physicist Richard Feynman (pronounced “FINE-men”). In addition to his groundbreaking work in theoretical physics (especially quantum electrodynamics and its associated Feynman diagrams), Feynman is known for The Feynman Lectures on Physics, which covers the elementary physics curriculum (mechanics, thermal physics, electrodynamics, etc.) in an unusually entertaining and insightful way. Let’s use square brackets and a boolean criterion on the "surname" column to find Feynman’s record in the laureates da.

What can you conclude from the format of this query below and results produce?

```
df[df['surname'] == 'Feynman']
```

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity |
|----|----|------------|---------|------------|------------|-------------|-----------------|-------------|
| 86 | 86 | Richard P. | Feynman | 1918-05-11 | 1988-02-15 | USA | US | New York NY |

The DataFrame object can take a Boolean condition on columns in it's index and will return records that meet that condition!. The inner part of the syntax for query above returns a Series consisting of boolean values for every laureate, with True if the surname is equal to "Feynman" and

False otherwise.

By using the correct index (i.e., 86), we can confirm that the value in that case is True.

Run the code cell below to confirm only this index returns True.

```
(df["surname"] == "Feynman") [86]

True
```

The `loc` attribute can be used in place of brackets in many places and is generally a more flexible way to pull out data items of interest. Let's use the `loc` attribute to retrieve the year when Feynman won.

Run the code cell below for this result.

```
df.loc[df["surname"] == "Feynman", "year"]

86    1965
Name: year, dtype: int64
```

Use the `loc` attribute illustrated above to code a query on this dataset to:

Find all Nobel laureates named named 'Curie'

Tip: Use the Boolean condition: `df["surname"].str.contains("Curie", na=False)`

```
# Code and test your scriplet to find all 'Curies'
# Assigned the results to a variable named curies.
curies = df.loc[df['surname'].str.contains("Curie", na=False)]
#View output
curies
```

| | id | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | |
|--|-----|-----------|---------|--------------|------------|-------------|-----------------------------|----------|--------|
| | 4 | 5 | Pierre | Curie | 1859-05-15 | 1906-04-19 | France | FR | Paris |
| | 5 | 6 | Marie | Curie | 1867-11-07 | 1934-07-04 | Russian Empire (now Poland) | PL | Warsaw |
| | 6 | 6 | Marie | Curie | 1867-11-07 | 1934-07-04 | Russian Empire (now Poland) | PL | Warsaw |
| | 191 | 194 | Irène | Joliot-Curie | 1897-09-12 | 1956-03-17 | France | FR | Paris |

Let's find all the winners of multiple nobel prizes using the `groupby` method

```
laureates = df.groupby(["id", "firstname", "surname"])
sizes = laureates.size()
sizes[sizes > 1] # result should show 5 winners

id  firstname  surname
6   Marie      Curie      2
66  John       Bardeen     2
217 Linus       Pauling     2
222 Frederick   Sanger     2
743 Barry      Sharpless  2
dtype: int64
```

Selecting Dates & Time information Pandas provides good support for datetimes. Let's search for laureates by exact birthday as a string - the default pandas storage format for dates.

Run the code cell below to get the result entry for Einstein born 3/14/1879 fondly known as Pi Day.

```
df.loc[df['born'] == '1879-03-14']
```

| id | | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity |
|----|----|-----------|----------|------------|------------|-------------|-----------------|----------|
| 25 | 26 | Albert | Einstein | 1879-03-14 | 1955-04-18 | Germany | DE | Ulm |

Like the query above let’s see if there are any laureates born 6/28 known as Tau Day

Run the code cell below to confirm!


```
df.loc[df['born'].str.contains('06-08', na=False)]
```

| id | | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity |
|-----|-----|------------|-----------|------------|------------|----------------|-----------------|---------------|
| 120 | 121 | Kenneth G. | Wilson | 1936-06-08 | 2013-06-15 | USA | US | Waltham |
| 370 | 372 | Francis | Crick | 1916-06-08 | 2004-07-28 | United Kingdom | GB | Northampton |
| 452 | 454 | Eric F. | Wieschaus | 1947-06-08 | 0000-00-00 | USA | US | South I |
| 667 | 683 | Wassily | Leontief | 1906-08-05 | 1999-02-05 | Russia | RU | Peters |
| 781 | 799 | Robert J. | Aumann | 1930-06-08 | 0000-00-00 | Germany | DE | Frankfur the- |
| 918 | 944 | Jacques | Dubochet | 1942-06-08 | 0000-00-00 | Switzerland | CH | . |

Rewrite the query above to filter the result to just laureates in Physics born on Tau Day ... Tip: use (df['category'] == "physics") as the other part of your Boolean function!

```
# Your code for the Physics category Tau Day Nobel Laureates below
tau_day_physics_curies = df.loc[
    (df['category'] == "physics") &
    (df['born'].str.contains('06-28', na=False))
]
```

tau_day_physics_curies



| id | | firstname | surname | born | died | bornCountry | bornCountryCode | bornCity | diedCountry | diedCountryCode | diedCity | g |
|-----|-----|-----------|----------------|------------|------------|-------------------------------------|-----------------|--------------------------|-------------|-----------------|--------------|---|
| 79 | 79 | Maria | Goeppert Mayer | 1906-06-28 | 1972-02-20 | Germany (now Poland) | PL | Kattowitz (now Katowice) | USA | US | San Diego CA | i |
| 125 | 126 | Klaus | von Klitzing | 1943-06-28 | 0000-00-00 | German-occupied Poland (now Poland) | PL | Schroda | NaN | NaN | NaN | |

